

NeuralNetworks

August 29, 2020

1 Neural Network training

```
[1]: import tensorflow as tf
from tensorflow import keras
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
```

```
[2]: (x_train, y_train), (x_test, y_test) = keras.datasets.mnist.load_data()

# Preprocess the data (these are NumPy arrays)
x_train = x_train / 255.0
x_test = x_test / 255.0

y_train = y_train.astype("float32")
y_test = y_test.astype("float32")

# Reserve 10,000 samples for validation
x_val = x_train[-10000:]
y_val = y_train[-10000:]
x_train = x_train[:-10000]
y_train = y_train[:-10000]
```

```
[8]: model = keras.models.Sequential()
model.add(keras.layers.Flatten(input_shape=[28, 28]))
model.add(keras.layers.Dense(300, activation="relu"))
model.add(keras.layers.Dense(300, activation="relu"))
model.add(keras.layers.Dense(10, activation="softmax"))
```

```
[13]: model.compile(loss="sparse_categorical_crossentropy",
                    optimizer="sgd",
                    metrics=["accuracy"])
```

```
[14]: early_stopping_cb = keras.callbacks.EarlyStopping(patience=5,
    ↳ restore_best_weights=True)
checkpoint_cb = keras.callbacks.ModelCheckpoint("classifier.h5",
    ↳ save_best_only=True)
```

```
history = model.fit(x_train, y_train, batch_size=128, epochs=100,  
                    validation_data=(x_val, y_val), callbacks=[checkpoint_cb, early_stopping_cb])
```

Epoch 1/100

391/391 [=====] - 1s 3ms/step - loss: 0.0171 -
accuracy: 0.9952 - val_loss: 0.0751 - val_accuracy: 0.9792

Epoch 2/100

391/391 [=====] - 1s 2ms/step - loss: 0.0142 -
accuracy: 0.9965 - val_loss: 0.0731 - val_accuracy: 0.9802

Epoch 3/100

391/391 [=====] - 1s 2ms/step - loss: 0.0128 -
accuracy: 0.9969 - val_loss: 0.0721 - val_accuracy: 0.9800

Epoch 4/100

391/391 [=====] - 1s 2ms/step - loss: 0.0120 -
accuracy: 0.9973 - val_loss: 0.0713 - val_accuracy: 0.9803

Epoch 5/100

391/391 [=====] - 1s 3ms/step - loss: 0.0113 -
accuracy: 0.9976 - val_loss: 0.0708 - val_accuracy: 0.9806

Epoch 6/100

391/391 [=====] - 1s 2ms/step - loss: 0.0109 -
accuracy: 0.9977 - val_loss: 0.0705 - val_accuracy: 0.9806

Epoch 7/100

391/391 [=====] - 1s 2ms/step - loss: 0.0104 -
accuracy: 0.9977 - val_loss: 0.0704 - val_accuracy: 0.9807

Epoch 8/100

391/391 [=====] - 1s 3ms/step - loss: 0.0101 -
accuracy: 0.9980 - val_loss: 0.0703 - val_accuracy: 0.9808

Epoch 9/100

391/391 [=====] - 1s 3ms/step - loss: 0.0098 -
accuracy: 0.9981 - val_loss: 0.0702 - val_accuracy: 0.9809

Epoch 10/100

391/391 [=====] - 1s 3ms/step - loss: 0.0095 -
accuracy: 0.9983 - val_loss: 0.0700 - val_accuracy: 0.9810

Epoch 11/100

391/391 [=====] - 1s 3ms/step - loss: 0.0093 -
accuracy: 0.9982 - val_loss: 0.0700 - val_accuracy: 0.9813

Epoch 12/100

391/391 [=====] - 1s 2ms/step - loss: 0.0090 -
accuracy: 0.9984 - val_loss: 0.0700 - val_accuracy: 0.9813

Epoch 13/100

391/391 [=====] - 1s 3ms/step - loss: 0.0088 -
accuracy: 0.9984 - val_loss: 0.0699 - val_accuracy: 0.9814

Epoch 14/100

391/391 [=====] - 1s 2ms/step - loss: 0.0086 -
accuracy: 0.9984 - val_loss: 0.0703 - val_accuracy: 0.9814

Epoch 15/100

391/391 [=====] - 1s 2ms/step - loss: 0.0085 -

```


accuracy: 0.9984 - val_loss: 0.0699 - val_accuracy: 0.9814
Epoch 16/100
391/391 [=====] - 1s 3ms/step - loss: 0.0083 -
accuracy: 0.9985 - val_loss: 0.0701 - val_accuracy: 0.9816
Epoch 17/100
391/391 [=====] - 1s 2ms/step - loss: 0.0081 -
accuracy: 0.9986 - val_loss: 0.0700 - val_accuracy: 0.9812
Epoch 18/100
391/391 [=====] - 1s 3ms/step - loss: 0.0080 -
accuracy: 0.9985 - val_loss: 0.0702 - val_accuracy: 0.9815
Epoch 19/100
391/391 [=====] - 1s 3ms/step - loss: 0.0078 -
accuracy: 0.9986 - val_loss: 0.0701 - val_accuracy: 0.9814
Epoch 20/100
391/391 [=====] - 1s 3ms/step - loss: 0.0077 -
accuracy: 0.9986 - val_loss: 0.0701 - val_accuracy: 0.9814

```

```

[15]: pd.DataFrame(history.history).plot(figsize=(8, 5))
plt.grid(True)
plt.gca().set_ylim(0, 1) # set the vertical range to [0-1]
plt.show()

```



```

[16]: model.evaluate(x_test, y_test)

```

```

313/313 [=====] - 0s 855us/step - loss: 0.0567 -

```

accuracy: 0.9833

[16]: [0.056724756956100464, 0.983299970626831]